

EKS Workshop v1.0

准备工作

IAM 配置

- EKS 集群角色 (通过 Console 创建时必须)
 - Trusted Entity: eks.amazonaws.com
 - AWS 托管策略:
 - AmazonEKSClusterPolicy
 - AmazonEKSServicePolicy
- EKS 节点组角色 (通过 Console 创建时必须)
 - Trusted Entity: ec2.amazonaws.com.cn
 - AWS 托管策略:
 - AmazonEKS_CNI_Policy
 - AmazonEKSWorkerNodePolicy
 - AmazonEC2ContainerRegistryReadOnly
- eksctl 管理用户 (通过命令行创建时必须)
 - AWS 托管策略:
 - AmazonEC2FullAccess
 - AWSCloudFormationFullAccess
 - AmazonEC2ContainerRegistryReadOnly
 - 客户托管策略:
 - 从以下 GitHub 拉取 IAM Policy

```
$ git clone https://github.com/ryanlao67/aws-iam-template.git
```

- 在 IAM 策略中依次创建

安装 EKS 相关工具

- 安装 eksctl: https://docs.aws.amazon.com/zh_cn/eks/latest/userguide/eksctl.html
 - 0.15.0 正式版本 :
 - Add support for KMS encryption provider (#1897)
 - Add support for China regions (#1860)
 - Add region Beijing (cn-north-1) (#1741)
 - Add region Ningxia (#1720)
 - Support addons for China regions, refactor setting container image for addons (#1867)
 - Add support for Kubernetes 1.15 (#1917 #1916)
 - Add support for Bottlerocket NodeGroups (#1918 #1919)

```
# MacOS: https://github.com/weaveworks/eksctl/releases/download/0.15.0/eksctl
# Linux: https://github.com/weaveworks/eksctl/releases/download/0.15.0/eksctl
# Windows: https://github.com/weaveworks/eksctl/releases/download/0.15.0/eksctl
```

```
# Example installation on MacOS
$ curl --silent --location "https://github.com/weaveworks/eksctl/releases/download/v0.15.0/eksctl_0.15.0_macos_arm64.tar.gz"
$ sudo mv /tmp/eksctl /usr/local/bin
$ eksctl version
0.15.0
```

- 安装 aws-iam-authenticator: https://docs.aws.amazon.com/zh_cn/eks/latest/userguide/install-aws-iam-authenticator.html
- 安装 Kubectl: https://docs.aws.amazon.com/zh_cn/eks/latest/userguide/install-kubectl.html
 - 中国区域可以直接从 kops-cn 的 GitHub 下载：

```
# Available version:
### v1.15.5 / v1.15.7
### v1.14.6 / v1.14.8
### v1.13.2 / v1.13.5 / v1.13.10 / v1.13.12
### v1.12.7 / v1.12.8 / v1.12.9 / v1.12.10
### v1.11.6 / v1.11.7 / v1.11.8 / v1.11.9
### v1.10.3 / v1.10.6 / v1.10.11
### v1.9.3 / v1.9.6 / v1.9.8
$ VERSION='<kubectL_version>'
$ curl -L https://s3.cn-north-1.amazonaws.com.cn/kops-bjs/fileRepository/kube
$ chmod +x kubectL
$ sudo mv kubectL /usr/local/bin
$ kubectL version --short --client
Client Version: <kubectL_version>
```

- 安装 AWS CLI: <https://docs.amazonaws.cn/cli/latest/userguide/install-cliv1.html>

准备密钥对

- 在 Console 上访问 EC2 页面，并定位到密钥对
- 点击 “Create key pair”，指定名称，例如：bjs
- 文件格式选择 pem，并点击创建密钥对
- 文件会下载至本地，打开 Terminal，并 cd 到密钥对下载目录
- 运行如下命令创建 pub 文件用于之后访问工作节点：

```
$ chmod 400 bjs.pem
$ ssh-keygen -y -f bjs.pem > bjs.pub
$ mv bjs.pub ~/.ssh/bjs.pub
```

准备 EKS 集群配置

CLUSTER YAML

```
$ vim cluster.yaml

apiVersion: eksctl.io/v1alpha5
kind: ClusterConfig

metadata:
  name: eks-demo-cluster
  region: cn-north-1
```

NODEGROUP1 YAML

- 创建一个由2个 m5.large 节点组成的 worker group，并部署在私有子网中
- 能够使用本地的 ssh key 登录相应的 worker 节点
- 制定两个label：
 - role: workers
 - usage: stable_cluster

```
$ vim nodegroup1.yaml

apiVersion: eksctl.io/v1alpha5
kind: ClusterConfig

metadata:
  name: eks-demo-cluster
  region: cn-north-1

nodeGroups:
  - name: eks-demo-workers1
    labels: { role: workers, usage: stable_cluster }
    instanceType: m5.large
    desiredCapacity: 2
    privateNetworking: yes
    ssh:
      publicKeyPath: ~/.ssh/bjs.pub
```

NODEGROUP2 YAML

- 创建一个由1个 c5.large 节点组成的 worker group，并部署在私有子网中
- 能够使用本地的 ssh key 登录相应的 worker 节点
- 制定两个label：
 - role: workers
 - usage: stable_cluster

```
$ vim nodegroup2.yaml

apiVersion: eksctl.io/v1alpha5
kind: ClusterConfig

metadata:
  name: eks-demo-cluster
  region: cn-north-1

nodeGroups:
  - name: eks-demo-workers2
    labels: { role: workers, usage: stable_cluster }
    instanceType: c5.large
    desiredCapacity: 1
    privateNetworking: yes
    ssh:
      publicKeyPath: ~/.ssh/bjs.pub
```

NODEGROUP3 YAML

- 创建一个由1个 t3.large 和 t3.medium 组成的 spot fleet worker group，并部署在私有子网中
- 可以根据实际用量进行弹性伸缩，最少2个节点，最多5个节点
- 能够使用本地的 ssh key 登录相应的 worker 节点
- 制定两个label：
 - role: workers
 - usage: spot_cluster

```
$ vim nodegroup3.yaml

apiVersion: eksctl.io/v1alpha5
kind: ClusterConfig

metadata:
  name: eks-demo-cluster
  region: cn-north-1

nodeGroups:
- name: eks-demo-workers3
  labels: { role: workers, usage: spot_cluster }
  minSize: 2
  maxSize: 5
  instancesDistribution:
    maxPrice: 0.2
    instanceTypes: ["t3.large", "t3.medium"]
    onDemandBaseCapacity: 0
    onDemandPercentageAboveBaseCapacity: 50
    spotInstancePools: 2
  privateNetworking: yes
  ssh:
    publicKeyPath: ~/.ssh/bjs.pub
```

创建 EKS 集群

创建 MASTER

```
# Create EKS cluster
$ eksctl create cluster \
  --config-file=cluster.yaml \
  --profile bjs

# Verify created cluster
$ eksctl get cluster \
  --profile bjs
NAME                      REGION
eks-demo-cluster         cn-north-1
```

创建 NODEGROUP

```
# Create node groups
$ eksctl create nodegroup \
  --config-file=nodegroup1.yaml \
  --profile bjs

$ eksctl create nodegroup \
  --config-file=nodegroup2.yaml \
  --profile bjs

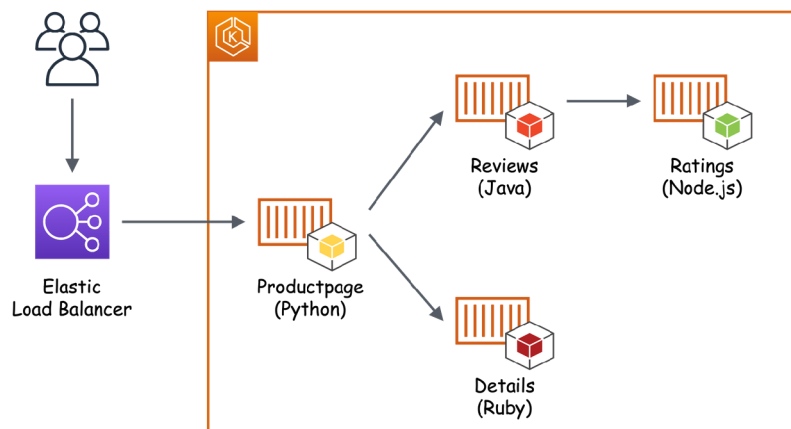
$ eksctl create nodegroup \
  --config-file=nodegroup3.yaml \
  --profile bjs

# Verify node groups
$ eksctl get nodegroup \
  --cluster eks-demo-cluster \
  --profile bjs
```

CLUSTER	NODEGROUP	CREATED	MIN SIZE	MAX SIZE	DESIF
eks-demo-cluster	eks-demo-workers1	2020-03-XXTXX:XX:XXZ	2	2	2
eks-demo-cluster	eks-demo-workers2	2020-03-XXTXX:XX:XXZ	1	1	1
eks-demo-cluster	eks-demo-workers3	2020-03-XXTXX:XX:XXZ	2	5	0

创建 Demo 应用

应用架构



制作镜像

- 从以下 GitHub 拉取示例应用代码
 - <https://github.com/ryanlao67/bookinfo-productpage.git>
 - <https://github.com/ryanlao67/bookinfo-reviews.git>
 - <https://github.com/ryanlao67/bookinfo-ratings.git>
 - <https://github.com/ryanlao67/bookinfo-details.git>

```
$ git clone https://github.com/ryanlao67/bookinfo-productpage.git
$ git clone https://github.com/ryanlao67/bookinfo-reviews.git
$ git clone https://github.com/ryanlao67/bookinfo-ratings.git
$ git clone https://github.com/ryanlao67/bookinfo-details.git
```

- 为 bookinfo-productpage 制作 Docker 镜像

```
$ cd bookinfo-productpage
$ docker build -t bookinfo-productpage:1.0 .
```

- 为 bookinfo-reviews 制作 Docker 镜像

```
# 原始代码需要先通过如下命令生成war包
$ cd bookinfo-reviews
$ docker run --rm -u root -v "$(pwd)":/home/gradle/project -w /home/gradle/project

# 生成war包后，制作Docker镜像
# 目前示例代码中已包含war包，可以通过如下命令直接制作Docker镜像
$ cd bookinfo-reviews/reviews-wlpcfg
$ docker build -t bookinfo-reviews:1.0 .
```

- 为 bookinfo-ratings 制作 Docker 镜像

```
$ cd bookinfo-ratings
$ docker build -t bookinfo-ratings:1.0 .
```

- 为 bookinfo-details 制作 Docker 镜像

```
$ cd bookinfo-details
$ docker build -t bookinfo-details:1.0 .
```

- 完成镜像制作后，查看本地镜像

```
$ docker images
```

REPOSITORY	TAG	IMAGE ID	CREATED
bookinfo-ratings	1.0	1ee994c63a8f	2 minutes ago
bookinfo-reviews	1.0	49772a269c04	3 minutes ago
bookinfo-details	1.0	09518b297a40	3 minutes ago
bookinfo-productpage	1.0	1cd26d55be02	13 minutes ago
python	3.6-slim	d3ae39a2a3a1	2 weeks ago
node	12-slim	26932a190e66	2 weeks ago
ruby	2.6.3-slim	1c75fac01180	7 months ago
websphere-liberty	19.0.0.5-javaee8	99f7cc7fc995	8 months ago
gradle	4.8.1	8b2989808a5c	20 months ago

镜像推送

- 通过 AWS CLI 创建 ECR 仓库

```
$ aws ecr create-repository \
  --repository-name eks-demo/bookinfo-productpage \
  --profile bjs

$ aws ecr create-repository \
  --repository-name eks-demo/bookinfo-reviews \
  --profile bjs

$ aws ecr create-repository \
  --repository-name eks-demo/bookinfo-ratings \
  --profile bjs
```

```
$ aws ecr create-repository \
  --repository-name eks-demo/bookinfo-details \
  --profile bjs
```

- 通过 AWS CLI 验证 ECR 仓库

```
$ aws ecr describe-repositories \
  --query 'repositories[*].[repositoryName]' \
  --output text \
  --profile bjs
# ECR显示如下
eks-demo/bookinfo-reviews
eks-demo/bookinfo-details
eks-demo/bookinfo-ratings
eks-demo/bookinfo-productpage
```

- 登录 ECR

```
$ $(aws ecr get-login --no-include-email --region cn-north-1 --profile bjs)
Login Succeeded
```

- 打标签并推送到创建的 ECR

```
# Tag images
# Account ID shown as 12-digits
$ docker tag bookinfo-productpage:1.0 855501529395.dkr.ecr.cn-north-1.amazonaws.c
$ docker tag bookinfo-reviews:1.0 855501529395.dkr.ecr.cn-north-1.amazonaws.com.c
$ docker tag bookinfo-ratings:1.0 855501529395.dkr.ecr.cn-north-1.amazonaws.com.c
$ docker tag bookinfo-details:1.0 855501529395.dkr.ecr.cn-north-1.amazonaws.com.c

# Push to ECR
# Account ID shown as 12-digits
$ docker push 855501529395.dkr.ecr.cn-north-1.amazonaws.com.cn/eks-demo/bookinfo-
$ docker push 855501529395.dkr.ecr.cn-north-1.amazonaws.com.cn/eks-demo/bookinfo-
$ docker push 855501529395.dkr.ecr.cn-north-1.amazonaws.com.cn/eks-demo/bookinfo-
$ docker push 855501529395.dkr.ecr.cn-north-1.amazonaws.com.cn/eks-demo/bookinfo-
```

创建应用 YAML

- bookinfo-productpage

```
$ vim bookinfo-productpage.yaml

apiVersion: v1
kind: Service
metadata:
  name: productpage
  labels:
    app: productpage
    service: productpage
spec:
  ports:
    - port: 9080
      name: http
  selector:
    app: productpage
```

```

    type: LoadBalancer
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: productpage-v1
  labels:
    app: productpage
    version: v1
spec:
  replicas: 2
  selector:
    matchLabels:
      app: productpage
      version: v1
  template:
    metadata:
      labels:
        app: productpage
        version: v1
    spec:
      serviceAccountName: default
      containers:
      - name: productpage
        image: 855501529395.dkr.ecr.cn-north-1.amazonaws.com.cn/eks-demo/bookinfo
        imagePullPolicy: IfNotPresent
        ports:
        - containerPort: 9080
      nodeSelector:
        usage: stable_cluster

```

- bookinfo-reviews

```
$ vim bookinfo-reviews.yaml
```

```

apiVersion: v1
kind: Service
metadata:
  name: reviews
  labels:
    app: reviews
    service: reviews
spec:
  ports:
  - port: 9080
    name: http
  selector:
    app: reviews
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: reviews-v1
  labels:
    app: reviews
    version: v1

```



```

spec:
  replicas: 1
  selector:
    matchLabels:
      app: reviews
      version: v1
  template:
    metadata:
      labels:
        app: reviews
        version: v1
    spec:
      serviceAccountName: default
      containers:
      - name: reviews
        image: 855501529395.dkr.ecr.cn-north-1.amazonaws.com.cn/eks-demo/bookinfo
        imagePullPolicy: Always
        ports:
        - containerPort: 9080
      affinity:
        podAffinity:
          requiredDuringSchedulingIgnoredDuringExecution:
          - labelSelector:
              matchExpressions:
              - key: app
                operator: In
                values:
                - productpage
            topologyKey: "kubernetes.io/hostname"
      nodeSelector:
        usage: stable_cluster

```

- bookinfo-ratings

```
$ vim bookinfo-ratings.yaml
```

```

apiVersion: v1
kind: Service
metadata:
  name: ratings
  labels:
    app: ratings
    service: ratings
spec:
  ports:
  - port: 9080
    name: http
  selector:
    app: ratings
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: ratings-v1
  labels:

```

```

    app: ratings
    version: v1
spec:
  replicas: 1
  selector:
    matchLabels:
      app: ratings
      version: v1
  template:
    metadata:
      labels:
        app: ratings
        version: v1
    spec:
      serviceAccountName: default
      containers:
      - name: ratings
        image: 855501529395.dkr.ecr.cn-north-1.amazonaws.com.cn/eks-demo/bookinfo
        imagePullPolicy: IfNotPresent
        ports:
        - containerPort: 9080
      nodeSelector:
        usage: spot_cluster

```

- bookinfo-details

```
$ vim bookinfo-details.yaml
```

```

apiVersion: v1
kind: Service
metadata:
  name: details
  labels:
    app: details
    service: details
spec:
  ports:
  - port: 9080
    name: http
  selector:
    app: details
---
apiVersion: apps/v1
kind: Deployment
metadata:
  name: details-v1
  labels:
    app: details
    version: v1
spec:
  replicas: 1
  selector:
    matchLabels:
      app: details
      version: v1

```

```

template:
  metadata:
    labels:
      app: details
      version: v1
  spec:
    serviceAccountName: default
    containers:
    - name: details
      image: 855501529395.dkr.ecr.cn-north-1.amazonaws.com.cn/eks-demo/bookinfo
      imagePullPolicy: IfNotPresent
      ports:
      - containerPort: 9080
    nodeSelector:
      usage: stable_cluster

```

部署 Demo 应用

- 通过 kubectl 部署上述4个服务

```

$ kubectl apply -f bookinfo-productpage.yaml
$ kubectl apply -f bookinfo-reviews.yaml
$ kubectl apply -f bookinfo-ratings.yaml
$ kubectl apply -f bookinfo-details.yaml

```

- 查看 pod 运行状态

```

$ kubectl get pods -owide

```

NAME	READY	STATUS	RESTARTS	AGE	IP
details-v1-56d6c64f9-jfxrc	1/1	Running	0	36s	192.168.180.1
productpage-v1-6cf5b9797-gjtzc	1/1	Running	0	61s	192.168.96.76
productpage-v1-6cf5b9797-lrdlw	1/1	Running	0	61s	192.168.174.2
ratings-v1-6d5885c596-sg2rv	1/1	Running	0	43s	192.168.103.2
reviews-v1-b75ddfc84-k2rzv	1/1	Running	0	52s	192.168.108.1

- 查看 service 运行状态

```

$ kubectl get svc

```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP
details	ClusterIP	10.100.76.194	<none>
kubernetes	ClusterIP	10.100.0.1	<none>
productpage	LoadBalancer	10.100.93.199	aee3c98f5645611eaa2d406ab5237157-91
ratings	ClusterIP	10.100.223.103	<none>
reviews	ClusterIP	10.100.240.8	<none>

- 访问应用 http://{elb_endpoint}:9080/productpage

参考

- eksctl 配置文件示例：<https://github.com/weaveworks/eksctl/tree/master/examples>
- 如果拉取镜像比较慢可以使用 nwcd 镜像，如下：

```
$ docker pull dockerhub.mirror.nwcdcdn.cn/library/<image>:<tag>
# 例如：docker pull dockerhub.mirror.nwcdcdn.cn/library/node:12-slim
$ docker tag dockerhub.mirror.nwcdcdn.cn/library/node:12-slim node:12-slim
```

思考

- Q1: 能不能通过一个 YAML 文件创建包含所有 NodeGroup 的 EKS 集群？
- Q2: 有什么方式可以快速实现应用之间的安全隔离吗？
- Q3: 怎么使用 ALB 作为运行在 EKS 集群里的应用的前端入口？